

DMML, 7 Mar 2019

Boosting a weak classifier

Start with a weak classifier h_0 over $T = \{w_1, \dots, w_N\}$

makes mistakes on T ,
accuracy $> \frac{1}{2}$

Assign weight $1/N$ to each input

Add up weight of mistakes — boost weight of
wrong inputs by a factor based on this

Build a classifier h_1 with new weights

Repeat

Sequence of weak classifiers h_0, h_1, \dots, h_m

Error rate e_0, e_1, \dots, e_m

Weighted sum of classifiers, using error rates

If classifiers vote as ± 1

$$h(x) = \text{sign}(\alpha_0 h_0(x) + \alpha_1 h_1(x) + \dots + \alpha_m h_m(x))$$

weight of
the classifier

$+1$ or
 -1

Related setting

- A collection of weak classifiers ("experts")
- Combine expert advice

Topic classifier for articles

Expert i - checks for some words
and gives a topic,
else says "no guess"

↑
"sleeping expert"

No sleeping - every expert classifies every input.

Experts: k_1, k_2, \dots

Find a subset $J \subseteq M$ experts and assign weights α_i to this subset

Chosen set: $\hat{k}_1, \hat{k}_2, \dots, \hat{k}_M$

Weights: $\alpha_1, \alpha_2, \dots, \alpha_M$

← need not be distinct

$$h(x) = \text{sign} \left(\sum_{i=1}^M \alpha_i \hat{k}_i(x) \right)$$

Weights for experts, weights for inputs
 α_i w_j $= w_j^{(0)}$
expert k_i input (x_j, y_j)

Initially $w_j = \frac{1}{N}$

Iteratively add $\hat{k}_1, \hat{k}_2, \dots, \hat{k}_m$ to our team
of experts

Inductively assume we have chosen

$\hat{k}_1, \dots, \hat{k}_m$ with weights $\alpha_1, \dots, \alpha_m$
input weights $w_j^{(m)}$

Iteration $m+1$

Run each k_i on training data and get weighted sum of errors

$$E_i = \sum_{y_j \neq k_i(x_j)} w_j^{(m)}$$

Pick \hat{k}_{m+1} as k_i with min E_i

$$e_{\min} = \frac{E_i}{W}$$

Sum of all weights

$$\alpha_m = \frac{1}{2} \ln \left(\frac{1 - e_{\min}}{e_{\min}} \right)$$

Update weights

$$w_j^{(m+1)} = w_j^{(m)} \cdot e^{\alpha_m} = w_j^{(m)} \sqrt{\frac{1 - e_m}{e_m}}$$

if new \hat{k}_{m+1} makes a mistake on x_j

$$= w_j^{(m)} e^{-\alpha_m} = w_j^{(m)} \sqrt{\frac{e_m}{1 - e_m}}$$

if new \hat{k}_{m+1} gets x_j correct

Sleeping Experts

No weights on inputs

Each expert k_i acts on a subset $T_i \subseteq T$

Assign weight $\alpha_i = 1$ to each k_i

Input $x_j \rightarrow K_j$ subset of experts active on x_j

$$W_{x_j} = \sum_{k_i \in K_j} \alpha_i$$
 — weighted sum of experts who offer opinion on x_j

Given $k_i \in K_j$, either $k_i(x_j) = y_j$ or

$m_i^{x_j} = 0$ — $k_i(x_j) \neq y_j$
 "miss" value
 $m_i^{x_j} = 1$

$$r_i^{x_j} = \left(\sum_{k_l \in K_j} \left(P_e^{x_j} \cdot m_e^{x_j} \right) \right)$$

$$P_e^{x_j} = \frac{\alpha_e}{w_{x_j}}$$

$(1 + \epsilon) - m_i^{x_j}$

Update

$$\alpha_i = \alpha_i \cdot (1 + \epsilon)^{r_i x_j}$$

How will this perform if the classification is skewed?

Minority case is "Yes" - most classifiers will have high accuracy overall?

Tune weight update to focus on minority case.

Neural networks

Perceptron

$$W \cdot x \geq b$$


Instead

$$W \cdot x + b \geq 0$$


↑
weight

↑
bias

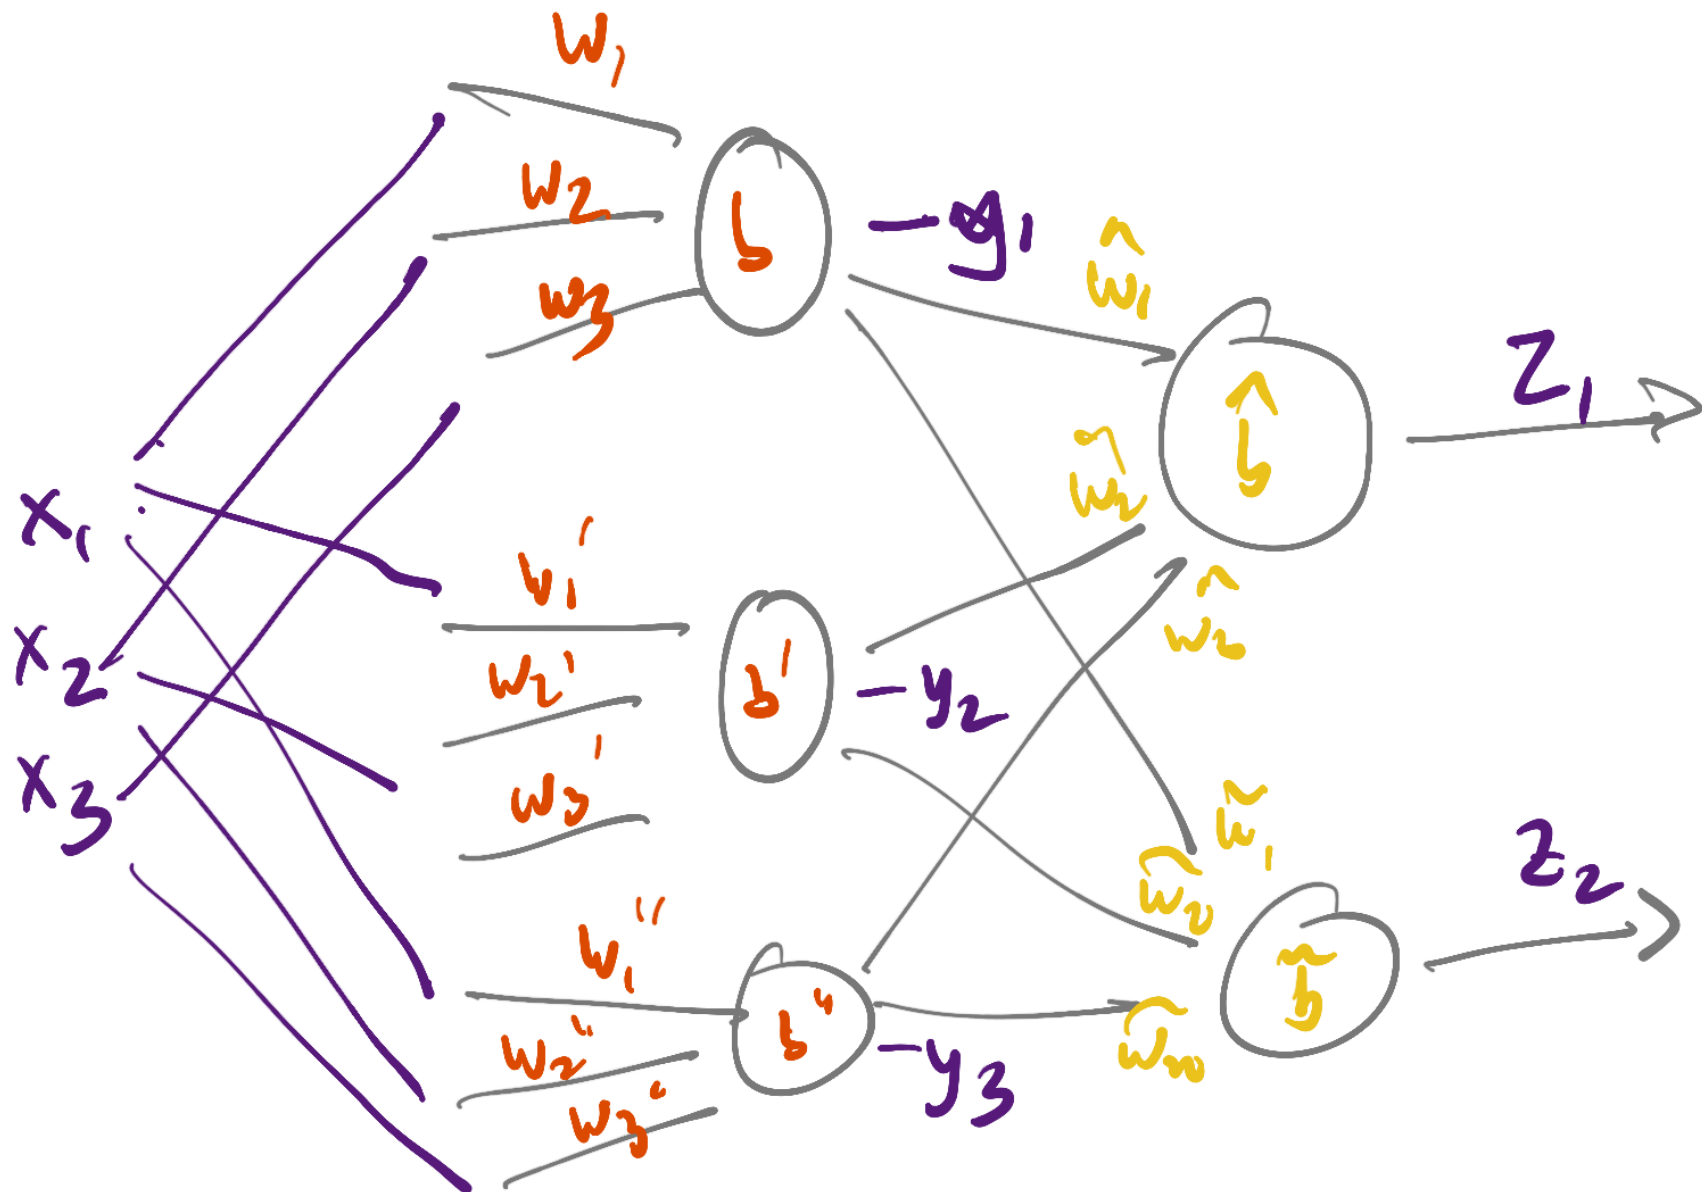
Pictorially




XOR function
cannot be modelled



Network of perceptrons





$$x_1, x_2 = 0, 1$$

NAND
gate

x_1	x_2	y
0	0	3
0	1	1
1	0	1
1	1	-1

$\left. \begin{matrix} 3 \\ 1 \\ 1 \end{matrix} \right\} \rightarrow 1$
 $\left. \begin{matrix} -1 \end{matrix} \right\} \rightarrow 0$

Functional completeness

OR + NOT can express any boolean function over (x_1, x_2)

AND + NOT

NAND

Perceptron for NAND \rightarrow networks of perceptrons are functionally complete